

# Lecture D1. Markov Reward Process 1

Sim, Min Kyu, Ph.D., [mksim@seoultech.ac.kr](mailto:mksim@seoultech.ac.kr)



서울과학기술대학교 데이터사이언스학과

- 1 I. Motivation
- 2 II. Method 1 - Monte-Carlo simulation
- 3 III. Method 2 - Iterative solution

# I. Motivation

# Recap

- In the first introduction of soda DTMC, the following question was posed.

*Given I drink coke today, what is likely my consumption for upcoming 10 days? (Pepsi is \$1 and Coke is \$1.5)*

- In Lecture note C1, Section 4, we demonstrated Monte-Carlo method that generates 10,000 (MC\_N) number of paths and found total expected cost to be approximately 13.36.
- This lecture builds more systematic approach rather than the previous time-consuming Monte-Carlo method.
- This lecture begins to introduce those daunting notations and mathematical treatment for reinforcement learning.

## reward and return

- Let the spending on day- $t$  is  $r_t$ . That is,  $r_t$  is *cost* or *reward* for time  $t$ .
- The reward  $r_t$  is determined by the state at time  $t$  with a function  $R(\cdot)$  such as  $r_t = R(s)$ .

### Definition 1 (reward function)

A real-valued function  $R : S \rightarrow \mathbb{R}$  is called a *reward function* that determines the reward given the state. That is,  $r_t = R(s)$ , where  $S_t = s$

- We were asked to find the expected value of  $r_0 + r_1 + \dots + r_9$ .

### Definition 2 (return)

The *return*  $G_t$  is the sum of remaining reward at time  $t$ .

- In our problem,
  - $G_0 = r_0 + r_1 + \dots + r_9$
  - $G_1 = r_1 + \dots + r_9$
  - $G_2 = r_2 + \dots + r_9$
  - ...
  - $G_9 = r_9$

- In our problem, we were asked to find the expected value of  $G_0$  starting from state  $c$  at time 0.
- At time 0, the value of  $r_0$  is known, but  $r_1, \dots, r_9$  are random variables. So,  $G_0$  is random variable as well.
- The random variable  $G_0$  depends on the current state  $S_0$  and the randomness along the stochastic path.
- In general, the random variable  $G_t$  depends on the last-known state  $S_t$  and some randomness along the remaining path.
- Since  $G_t$  is a random variable, we want to evaluate  $\mathbb{E}[G_t]$ . In particular, considering its dependence structure, we are interested in evaluating  $\mathbb{E}[G_t | S_t = s]$ .
- In this light, the current problem is  $\mathbb{E}[G_0 | S_0 = c]$ , or  $\mathbb{E}[r_0 + r_1 + \dots + r_9 | S_0 = c]$ .
- This motivates the following definition.

## state-value function

### Definition 3 (state-value function)

A *state-value function*  $V_t(s)$  is the expected return given state  $s$  at time  $t$ . That is,  
$$V_t(s) = \mathbb{E}[G_t | S_t = s].$$

- Again, we are interested in finding

$$V_0(c) = \mathbb{E}[G_0 | S_0 = c] = \mathbb{E}[r_0 + \dots + r_9 | S_0 = c].$$



## II. Method 1 - Monte-Carlo simulation

# Recap

- The MC simulation is a valid approach. We shall review our initial effort with newly introduced terminology.
- The algorithm includes...
  - 1 Generate a single stochastic path starting from the *initial state*,  $S_0 = c$ .
  - 2 Collect a single value of *return*,  $G_i$ ,  $1 \leq i \leq MC\_N$ , by accumulating *rewards*,  $\{r_0, r_1, \dots, r_9\}$ , along the path.
  - 3 Take an average of collected *returns* to evaluate state-value function,  $V_0(c)$ .

```
MC_N <- 10000
spending_records <- rep(0, MC_N)
for (i in 1:MC_N) {
  path <- "c" # coke today (day-0)
  for (t in 1:9) {
    this_state <- str_sub(path, -1, -1)
    next_state <- soda_simul(this_state)
    path <- paste0(path, next_state)
  }
  spending_records[i] <- cost_eval(path)
}

cost_eval <- function(path) {
  cost_one_path <-
    str_count(path, pattern = "c")*1.5 +
    str_count(path, pattern = "p")*1
  return(cost_one_path)
}
```

## MC simulation for estimating *state-value function*

- Formally, for a *finite-horizon MRP*, the following is MC simulation for estimating *state-value function*.

```
# MC evaluation for state-value function
# with state s, time 0, reward r, time-horizon H
1: episode_i <- 0
2: cum_sum_G_i <- 0
3: while episode_i < num_episode
4:   Generate an stochastic path starting from state s and time 0.
5:   Calculate return G_i <- sum of rewards from time 0 to time H-1.
6:   cum_sum_G_i <- cum_sum_G_i + G_i
7:   episode_i <- episode_i + 1
8: State-value-fn V_t(s) <- cum_sum_G_i/num_episode
9: return V_t(s)
```

### III. Method 2 - Iterative solution

# Motivation

- Same as previous section, our goal is still to estimate  $V_0(c) = \mathbb{E}[G_0|S_t = c]$ .
- Since  $G_t = \sum_{i=t}^9 r_i$  has less number of terms when  $t$  is high number, we shall start from  $t = 9$  and work backward, i.e. from  $V_9(s)$ , then  $V_8(s)$ , then  $V_7(s)$ , ...
- For  $t = 9$ ,
  - From the general formula  $V_t(s) = \mathbb{E}[G_t|S_t = s]$ , it is easy to see that  $V_9(s) = \mathbb{E}[G_9|S_9 = s] = \mathbb{E}[\sum_{i=9}^9 r_i|S_9 = s] = \mathbb{E}[r_9|S_9 = s] = R(s)$ .
  - In other words,
    - $V_9(c) = \mathbb{E}[r_9|S_9 = c] = R(c) = 1.0$  and
    - $V_9(p) = \mathbb{E}[r_9|S_9 = p] = R(p) = 1.5$ .
  - In general,

$$V_9(s) = R(s) + V_{10}(s) \tag{1}$$

, where  $V_{10}(s) = 0, \forall s$

- For  $t = 8$ ,
  - From the general formula  $V_t(s) = \mathbb{E}[G_t | S_t = s]$ , (watch below carefully)

$$\begin{aligned}
 V_8(s) &= \mathbb{E}[G_8 | S_8 = s] \\
 &= \mathbb{E}\left[\sum_{i=8}^9 r_i \mid S_8 = s\right] \\
 &= \mathbb{E}[r_8 + r_9 | S_8 = s] + \\
 &= \mathbb{E}[r_8 | S_8 = s] + \mathbb{E}[r_9 | S_8 = s] \\
 &= R(s) + \mathbb{E}[r_9 | S_8 = s]
 \end{aligned} \tag{2}$$

- Here, let's consider  $\mathbb{E}[r_9 | S_8 = c]$  first.
  - This is expected spending on day-9 given that I drink coke on day-8. This value is conditioned on what I drink on day-9. If coke on day-9 with probability 0.7,  $r_9 = 1.5$ . If pepsi w/ prob. 0.3,  $r_9 = 1.0$ . This expectation is 1.35 ( $= 0.7 \cdot 1.5 + 0.3 \cdot 1.0$ ).
  - Formally,

$$\begin{aligned}
 \mathbb{E}[r_9 | S_8 = c] &= \mathbf{P}_{cc} \mathbb{E}[r_9 | S_8 = c, S_9 = c] + \mathbf{P}_{cp} \mathbb{E}[r_9 | S_8 = c, S_9 = p] \\
 &= \mathbf{P}_{cc} \mathbb{E}[r_9 | S_9 = c] + \mathbf{P}_{cp} \mathbb{E}[r_9 | S_9 = p] \quad (\because \text{Markov property}) \\
 &= \mathbf{P}_{cc} \mathbb{E}[G_9 | S_9 = c] + \mathbf{P}_{cp} \mathbb{E}[G_9 | S_9 = p] \\
 &= \mathbf{P}_{cc} V_9(c) + \mathbf{P}_{cp} V_9(p)
 \end{aligned}$$

• (Cont'd for  $t = 8$ )

- Now, let's consider  $\mathbb{E}[r_9 | S_8 = s]$  for generalized state  $s$ . With the notation assuming a transition from this state  $s$  to the next state  $s'$ ,

$$\begin{aligned}
 \mathbb{E}[r_9 | S_8 = s] &= \sum_{s' \in S} \mathbf{P}_{ss'} \mathbb{E}[r_9 | S_8 = s, S_9 = s'] \\
 &= \sum_{s' \in S} \mathbf{P}_{ss'} \mathbb{E}[r_9 | S_9 = s'] \quad (\because \text{Markov property}) \\
 &= \sum_{s' \in S} \mathbf{P}_{ss'} \mathbb{E}[G_9 | S_9 = s'] \quad (\because \text{Markov property}) \\
 &= \sum_{s' \in S} \mathbf{P}_{ss'} V_9(s')
 \end{aligned} \tag{3}$$

- We shall now summarize for  $t = 8$ ,

$$\begin{aligned}
 V_8(s) &= \mathbb{E}[G_8 | S_8 = s] = \mathbb{E}[r_8 + G_9 | S_8 = s] \\
 &= R(s) + \mathbb{E}[G_9 | S_8 = s] \\
 &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_9(s')
 \end{aligned} \tag{4}$$

$$(\text{expected return at time 8}) = (\text{reward at time 9}) + (\text{expected return at time 9})$$

- For  $t = 7$ ,

- From the general formula  $V_t(s) = \mathbb{E}[G_t | S_t = s]$ ,

$$\begin{aligned}
 V_7(s) &= \mathbb{E}[G_7 | S_7 = s] \\
 &= \mathbb{E}\left[\sum_{i=7}^9 r_i \mid S_7 = s\right] \\
 &= \mathbb{E}[r_7 + r_8 + r_9 | S_7 = s] \\
 &= \mathbb{E}[r_7 | S_7 = s] + \mathbb{E}[r_8 + r_9 | S_7 = s] \\
 &= R(s) + \mathbb{E}[G_8 | S_7 = s]
 \end{aligned} \tag{5}$$

- You get the hint? From here, we want to use  $V_8(s) = \mathbb{E}[G_8 | S_8 = s]$  to express this as a recursive formula for state-value function just like Eq. (4).

$$\begin{aligned}
 V_7(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} \mathbb{E}[G_8 | S_7 = s, S_8 = s'] \\
 &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_8(s')
 \end{aligned} \tag{6}$$



- For general  $t$ , (*exercise*)

- So far,

$$V_{10}(s) = 0$$

$$V_9(s) = R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_{10}(s') \text{ from Eq. (1)}$$

$$V_8(s) = R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_9(s') \text{ from Eq. (4)}$$

$$V_7(s) = R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_8(s') \text{ from Eq. (6)}$$

$$\dots = \dots$$

$$V_t(s) = R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_{t+1}(s')$$

$$\dots = \dots$$

$$V_0(s) = R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_1(s')$$

- Note that the array of equations can be solve from the top to the bottom.
- This iterative method is called as *backward induction* that works well with finite horizon problem.
- This iterative method (and its painful derivaion) is the most important mathematical essence of Markov decision process.

# Implementation strategy

- Summary so far

$$\begin{aligned}
 V_{10}(s) &= 0 \\
 V_t(s) &= R(s) + \sum_{s' \in S} \mathbf{P}_{ss'} V_{t+1}(s') \quad (\text{for } t \in \{0, 1, \dots, 9\})
 \end{aligned}$$

- Strategy

- Column vector  $v_t$  for  $V_t(s)$
- Column vector  $R$  for  $R(s)$
- The term  $\sum_{s' \in S} \mathbf{P}_{ss'} V_{t+1}(s')$  can be written as  $\mathbf{P}v_{t+1}$ .
- It follows  $v_t = R + \mathbf{P}v_{t+1}$ .

```

P <- array(c(0.7,0.5,0.3,0.5), dim=c(2,2))
R <- array(c(1.5,1.0), dim=c(2,1))
H <- 10 # time-horizon
v_t1 <- array(c(0,0), dim=c(2,1)) # v_{t+1}

t <- H-1
while (t >= 0) {
  v_t <- R + P %*% v_t1
  t <- t-1
  v_t1 <- v_t
}
v_t

##           [,1]
## [1,] 13.35937
## [2,] 12.73438

```

- Thus, we have the following *state-value function*.

- $V_0(c) = 13.359375$
- $V_0(p) = 12.734375$

## Backward induction for estimating *state-value* function

- Formally, for a *finite-horizon MRP*, the following is *backward induction* for estimating *state-value* function.

```
# Backward induction for state-value function
# with transition prob mat P, reward vector R, time-horizon H, state-value vector v_{t}
1: v_H <- zero-column vector
2: t <- H-1
3: while t >= 0
4:   v_t <- R + P*v_{t+1}
5:   t <- t-1
9: return v_t # this is v_0(s) for all s, because t=0 at this point
```

```
cat(str)
```

```
## If I only had an hour to chop down a tree, I would spend the first 45 minutes sharpening my axe. -  
A. Lincoln
```