

## Lecture D2. Markov Reward Process 2

Sim, Min Kyu, Ph.D., [mksim@seoultech.ac.kr](mailto:mksim@seoultech.ac.kr)



서울과학기술대학교 데이터사이언스학과

- 1 I. Motivation
- 2 II. Method 3 - Analytic solution
- 3 III. Method 4 - Iterative solution - by fixed point theorem

# I. Motivation

## Recap

- A Markov chain is a stochastic process with the specification of
  - a state space  $S$
  - a transition probability matrix  $\mathbf{P}$
- A Markov reward process is a Markov chain with the specification of
  - a reward  $r_t$  with the reward function  $R(s)$
  - a time horizon  $H$ , which is the duration we are interested in cumulative sum of rewards.
- If  $H$  is finite, then we call *finite-horizon MRP*.
- If  $H$  is infinite, then we call *infinite-horizon MRP*.

## Formulating an infinite horizon MRP

- In the previous lecture, we dealt with the following question.

*Given I drink coke today, what is likely my consumption for upcoming 10 days? (Pepsi is \$1 and Coke is \$1.5)*

- Infinite horizon problem is such as following.

*I am to live eternally. Given I drink coke today, what is likely my consumption for my upcoming forever life? (Pepsi is \$1 and Coke is \$1.5)*

- It may seem unrealistic on this soda problem to have an infinite time horizon. But infinite horizon model is indeed more common for MRP due to following reasons.
- ① Time horizon may be finite, but the horizon may be believed to be a long time and/or  $H$  is not certain.
- ② In accounting principle, all businesses are assumed to be perpetual.
- ③ Really long finite time horizon can be approximated by infinite time.
- ④ Oftentimes, each time step is very small such as minute ,or even millisecond, making the number of total time step as a very large number.

## Return for infinite horizon

- Return for finite horizon in the previous lecture was

$$G_t = \sum_{i=t}^{H-1} r_i$$

- If extended for infinite horizon, it becomes

$$G_t = \sum_{i=t}^{\infty} r_i$$

- Even if  $r_i$  is a small number, the quantity is diverging as long as  $r_i$  does not decay drastically.
- cf)  $\sum 1/n = \infty$  and  $\sum 1/n^2 < \infty$
- In case  $r_i$  decaying drastically, the convergence is only guaranteed if the chain will eventually be absorbed to one of absorbing states whose rewards are zero.

## Discount factor

- A mathematically convenient way to guarantee is to introduce *discount factor*,  $\gamma < 1$
- Using a discount factor, the return becomes

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

- Or, it can be written as

$$G_t = \sum_{i=t}^{\infty} \gamma^{i-t} r_i$$

- Note that this generalizes the previous notation with  $\gamma = 1$



- Other than for computational reason for return convergence, many real problems indeed should be modelled with discount factor.
- Humans behave in much the same way, putting more importance in the near future.
- Interest rate is generally positive, making today's money worth more than tomorrow's money.
- Future is risky to some degree, making future's reward less valuable than today's reward.
- ~~If you die today, there is no tomorrow.~~

## State-value function

- Like before, the state-value function  $V_t(s)$  for a MDP and a state  $s$  is defined as the expected return starting from state  $s$  at time  $t$ , namely,

$$V_t(s) = \mathbb{E}[G_t | S_t = s]$$

- For infinite horizon problem, are the following two quantity different?
  - ①  $V_0(s) = \mathbb{E}[G_0 | S_t = 0]$
  - ②  $V_t(s) = \mathbb{E}[G_t | S_t = s]$
- It is not! This makes our life easier, and allowing us to drop the time subscript for the state-value function when necessary. Namely,  $V_0(s) = V_t(s) = V(s)$ .

# Summary

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots \quad (1)$$

$$G_t = \sum_{i=t}^{\infty} \gamma^{i-t} r_i \quad (2)$$

$$V_t(s) = \mathbb{E}[G_t | S_t = s] \quad (3)$$

## II. Method 3 - Analytic solution

## Development

- For a finite horizon MRP, the goal was to find  $V_t(s)$  for all states  $s$  for  $0 \leq t \leq H$ .
- Since  $V_0(s) = V_t(s) = V(s)$ , the goal is only to find  $V(s)$  for all states  $s$ .

$$\begin{aligned} V(s) &= V_t(s) = \mathbb{E}[G_t | S_t = s] \\ &= \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots | S_t = s] \\ &= R(s) + \gamma \mathbb{E}[r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots | S_t = s] \\ &= R(s) + \gamma \mathbb{E}[G_{t+1} | S_t = s] \\ &= R(s) + \gamma \sum_{\forall s'} \mathbb{P}[S_{t+1} = s' | S_t = s] \mathbb{E}[G_{t+1} | S_t = s, S_{t+1} = s'] \\ &= R(s) + \gamma \sum_{\forall s'} \mathbf{P}_{ss'} \mathbb{E}[G_{t+1} | S_{t+1} = s'] \\ &= R(s) + \gamma \sum_{\forall s'} \mathbf{P}_{ss'} V_{t+1}(s') \\ &= R(s) + \gamma \sum_{\forall s'} \mathbf{P}_{ss'} V(s') \end{aligned} \tag{4}$$

- The section for Iterative solution in the previous lecture had the last equation of

$$V_t(s) = R(s) + \sum_{\forall s'} \mathbf{P}_{ss'} V_{t+1}(s')$$

- The Eq (4) was

$$V(s) = R(s) + \gamma \sum_{\forall s'} \mathbf{P}_{ss'} V(s')$$

- These are very similar except that the previous one had  $\gamma = 1$ .
- It is constructed as

(Expected return at time  $t$ ) = (reward at time  $t$ ) + (Expected return at time  $t+1$ )

- These are called *Bellman's equation*, named after Richard R. Bellman ([wiki link](#)) who introduced dynamic programming in 1953.

# Analytic formula

$$V(s) = R(s) + \gamma \sum_{\forall s'} \mathbf{P}_{ss'} V(s')$$

- Once again, the strategy is
  - Column vector  $v$  for  $V(s)$
  - Column vector  $R$  for  $R(s)$
  - $\gamma \mathbf{P}v$  for  $\gamma \sum_{\forall s'} \mathbf{P}_{ss'} V(s')$
  - (where  $\mathbf{P}$  is a transition matrix)
  - It follows  $v = R + \gamma \mathbf{P}v$
- This can be solved as:

$$v = R + \gamma \mathbf{P}v$$

$$\Rightarrow Iv = R + \gamma \mathbf{P}v$$

$$\Rightarrow Iv - \gamma \mathbf{P}v = R$$

$$\Rightarrow (I - \gamma \mathbf{P})v = R$$

$$\Rightarrow v = (I - \gamma \mathbf{P})^{-1} R$$

## Example

*I am to live eternally. Given I drink coke today, what is likely my consumption for my upcoming forever life? (Pepsi is \$1 and Coke is \$1.5)*

- We need information regarding the discount rate. Let's assume  $\gamma = 0.9$ .
- We have

$$\begin{aligned}
 v &= R + \gamma \mathbf{P}v \\
 \begin{pmatrix} v(c) \\ v(p) \end{pmatrix} &= \begin{pmatrix} R(c) \\ R(p) \end{pmatrix} + \gamma \begin{pmatrix} \mathbf{P}_{cc} & \mathbf{P}_{cp} \\ \mathbf{P}_{pc} & \mathbf{P}_{pp} \end{pmatrix} \begin{pmatrix} v(c) \\ v(p) \end{pmatrix} \\
 \begin{pmatrix} v(c) \\ v(p) \end{pmatrix} &= \begin{pmatrix} 1.5 \\ 1.0 \end{pmatrix} + 0.9 \begin{pmatrix} 0.7 & 0.3 \\ 0.5 & 0.5 \end{pmatrix} \begin{pmatrix} v(c) \\ v(p) \end{pmatrix} \tag{5}
 \end{aligned}$$



```

P <- array(c(0.7,0.5,0.3,0.5), dim=c(2,2))
R <- array(c(1.5,1.0), dim=c(2,1))
gamma = .9
v <- solve(diag(2)-gamma*P)%*%R #  $v = (I - \gamma P)^{-1} R$ 
v

##           [,1]
## [1,] 13.35366
## [2,] 12.74390

```

## Exercise 1

*What is the relationship between the above vector  $v$  and stationary distribution?*

## Exercise 2

*What are your concerns for this approach?*

### III. Method 4 - Iterative solution - by fixed point theorem

# Recap

- The previous approach was based on the following two formula

$$v = R + \gamma \mathbf{P}v \quad (6)$$

$$v = (I - \gamma \mathbf{P})^{-1} R \quad (7)$$

- The Eq. (6) is a Bellman's equation.
- The Eq. (7) is used to find a analytic solution.
- Using the Eq (7), there are two concerns that you should have. (This is the suggested solution to Exercise 2)
  - 1 The matrix  $I - \gamma \mathbf{P}$  may not be invertible.
  - 2 Even if it's invertible, it may be prohibitive if the size of the matrix is big.
- We are free from the first concern. The matrix  $I - \gamma \mathbf{P}$  can be proved to be invertible always.
- We are not free from the second concern. So, this section introduces an alternative, numerical, and iterative approach.

# Iterative algorithm

- Using the fixed-point theorem along with Eq. (6), we apply the following iterative algorithm to find  $v$ .

$$v_{i+1} \leftarrow R + \gamma \mathbf{P} v_i$$

```

1: Let epsilon <- 10^{-8} # or some small number
2: Let v_0 <- zero vector
3: Let v_1 <- R + \gamma * P * v_0
4: i <- 1
5: While ||v_i - v_{i-1}|| > epsilon # may use any norm
6:   v_{i+1} <- R + \gamma * P * v_i
7:   i <- i+1
8: Return v_{i+1}

```

# Math Review - Norm

## Definition 1

For a length- $n$  vector  $x$ , the norm of vector  $\|x\|_p$  is defined as follows.

- 1-norm:  $\|x\|_1 = \sum_{i=1}^n |x_i|$  (*sum of absolute value*)
- 2-norm:  $\|x\|_2 = (\sum_{i=1}^n x_i^2)^{1/2}$  (*Euclidean distance, distance from the origin*)
- $\infty$ -norm:  $\|x\|_\infty = \max_{1 \leq i \leq n} |x_i|$  (*farthest axis*)

- Throughout this course, we will use  $\infty$ -norm to guarantee that value functions (or any other quantities) are well approximated for every state.

# Implementation

## ● The psedo code

```

1: Let  $\epsilon \leftarrow 10^{-8}$  # or some small number
2: Let  $v_0$  <- zero vector
3: Let  $v_1 \leftarrow R + \gamma P v_0$ 
4:  $i \leftarrow 1$ 
5: While  $\|v_i - v_{i-1}\| > \epsilon$  # may use any norm
6:    $v_{i+1} \leftarrow R + \gamma P v_i$ 
7:    $i \leftarrow i+1$ 
8: Return  $v_{i+1}$ 

```

## ● The R-code

● (I wish there was a do-while loop in R)

```

R <- array(c(1.5,1.0), dim=c(2,1))
P <- array(c(0.7,0.5,0.3,0.5), dim=c(2,2))
gamma <- 0.9
epsilon <- 10^(-8)
v_old <- array(rep(0,2), dim=c(2,1))
v_new <- R + gamma*P%*%v_old
while (max(abs(v_new-v_old)) > epsilon) { # inf-norm
  v_old <- v_new
  v_new <- R + gamma*P%*%v_old
}
print(v_new)

##           [,1]
## [1,] 13.35366
## [2,] 12.74390

```

## ● The full iteration process

```
R <- array(c(1.5,1.0), dim=c(2,1))
P <- array(c(0.7,0.5,0.3,0.5), dim=c(2,2))
gamma <- 0.9
epsilon <- 10^(-8)
v_old <- array(rep(0,2), dim=c(2,1))
v_new <- R + gamma*P%*%v_old
results <- t(v_old) # to save
results <- rbind(results, t(v_new)) # to save
while (max(abs(v_new-v_old)) > epsilon) {
  v_old <- v_new
  v_new <- R + gamma*P%*%v_old
  results <- rbind(results, t(v_new)) # to save
}
```

```
results <- data.frame(results)
colnames(results) <- c("coke", "pepsi")
```


```
head(results)
```

```
##      coke   pepsi
## 1 0.000000 0.000000
## 2 1.500000 1.000000
## 3 2.715000 2.125000
## 4 3.784200 3.178000
## 5 4.742106 4.132990
## 6 5.603434 4.993793
```

```
tail(results)
```

```
##      coke   pepsi
## 175 13.35366 12.7439
## 176 13.35366 12.7439
## 177 13.35366 12.7439
## 178 13.35366 12.7439
## 179 13.35366 12.7439
## 180 13.35366 12.7439
```






- The previous plot was generated by the following code.


```
library(tidyverse)
results$idx <- as.numeric(row.names(results))
results <- results %>%
  gather("coke", "pepsi", key="soda", value="value_fn")
ggplot(results, aes(x=idx, y=value_fn, group = soda, color = soda)) +
  geom_point() +
  theme_minimal()
```

- Note that there are quite convergence going on after many steps.
- After 50 steps (coke only)
- After 100 steps (coke only)

```
results %>% filter(idx >= 50, soda == "coke") %>%
  ggplot(aes(x=idx, y=value_fn)) +
  geom_point() +
  theme_minimal()
```



```
results %>% filter(idx >= 100, soda == "coke") %>%
  ggplot(aes(x=idx, y=value_fn)) +
  geom_point() +
  theme_minimal()
```



"Success isn't permanent, and failure isn't fatal. - Mike Ditka"