

## Lecture E2. MDP with Model 2

Sim, Min Kyu, Ph.D., [mksim@seoultech.ac.kr](mailto:mksim@seoultech.ac.kr)



서울과학기술대학교 데이터사이언스학과

1 I. Recap

2 II. Policy improvement

3 III. Policy iteration

## I. Recap

## policy\_eval()

```
gamma <- 1
states <- as.character(seq(0, 70, 10))
P_normal <- matrix(c(0,1,0,0,0,0,0,0,
                    0,0,1,0,0,0,0,0,
                    0,0,0,1,0,0,0,0,
                    0,0,0,0,1,0,0,0,
                    0,0,0,0,0,1,0,0,
                    0,0,0,0,0,0,1,0,
                    0,0,0,0,0,0,0,1,
                    0,0,0,0,0,0,0,0,1),
                  nrow = 8, ncol = 8, byrow = TRUE,
                  dimnames = list(states, states))
P_speed <- matrix(c(.1, 0,.9, 0, 0, 0, 0, 0,
                   .1, 0, 0,.9, 0, 0, 0, 0,
                   0,.1, 0, 0,.9, 0, 0, 0,
                   0, 0,.1, 0, 0,.9, 0, 0,
                   0, 0, 0,.1, 0, 0,.9, 0,
                   0, 0, 0, 0,.1, 0, 0,.9,
                   0, 0, 0, 0, 0,.1, 0,.9,
                   0, 0, 0, 0, 0, 0, 0, 1),
                 nrow = 8, ncol = 8, byrow = TRUE,
                 dimnames = list(states, states))
```

```
transition <- function(given_pi,
                      states, P_normal, P_speed) {
  P_out <- array(0,
                dim = c(length(states), length(states)),
                dimnames = list(states, states))
  for (s in states) {
    action_dist <- given_pi[s,]
    P <- action_dist["normal"]*P_normal +
          action_dist["speed"]*P_speed
    P_out[s,] <- P[s,]
  }
  return(P_out)
}
R_s_a <- matrix(
  c( -1,  -1,  -1,  -1,  0.0,  -1,  -1,  0,
     -1.5,-1.5,-1.5,-1.5,-0.5,-1.5,-1.5, 0),
  nrow = length(states), ncol = 2, byrow = FALSE,
  dimnames = list(states, c("normal", "speed")))
reward_fn <- function(given_pi, R_s_a) {
  R_pi <- rowSums(given_pi*R_s_a)
  return(R_pi)
}
```

```

policy_eval <- function(given_pi) {
  R <- reward_fn(given_pi, R_s_a = R_s_a)
  P <- transition(given_pi, states = states, P_normal = P_normal, P_speed = P_speed)
  gamma <- 1.0
  epsilon <- 10^(-8)
  v_old <- array(rep(0,8), dim=c(8,1))
  v_new <- R + gamma*P%%v_old
  while (max(abs(v_new-v_old)) > epsilon) {
    v_old <- v_new
    v_new <- R + gamma*P%%v_old
  }
  return(v_new)
}

pi_speed <- cbind(rep(0,length(states)), rep(1,length(states)))
rownames(pi_speed) <- states; colnames(pi_speed) <- c("normal", "speed")
t(policy_eval(pi_speed))

```

```

##           0           10           20           30           40           50           60 70
## [1,] -5.805929 -5.208781 -4.139262 -3.475765 -2.35376 -1.735376 -1.673538 0

```

```

pi_50 <- cbind(rep(0.5,length(states)), rep(0.5,length(states)))
rownames(pi_50) <- states; colnames(pi_50) <- c("normal", "speed")
t(policy_eval(pi_50))

```

```

##           0           10           20           30           40           50           60 70
## [1,] -5.969238 -5.133592 -4.119955 -3.389228 -2.04147 -2.027768 -1.351388 0

```

## Major components of approaching MDP

- ❶ **(policy evaluation)** We need to be able to evaluate  $V^\pi(s)$  for a fixed  $\pi$ . This is called *policy evaluation*. This is also called as *prediction* in reinforcement learning.
  - ❷ **(optimal value function)** We want to be able to evaluate  $V^{\pi^*}(s)$ , where  $\pi^*$  is the *optimal policy*. The quantity,  $V^{\pi^*}(s)$ , is optimal policy's value function, or called shortly as *optimal value function*.
  - ❸ **(optimal policy)** We want to find the *optimal policy*  $\pi^*$ . This is also called as *control* in reinforcement learning
- Check your reasoning why the followings are possible.
    - Optimal policy first: (optimal policy) + (policy evaluation)  $\rightarrow$  (optimal value function)
    - Optimal value function first: (optimal value function)  $\rightarrow$  (optimal policy)
  - This note will discuss
    - (policy evaluation) + *series of (policy improvement)*  $\rightarrow$  (optimal policy)

## II. Policy improvement

## Development

- Remind that we have a Bellman's equation for MDP as follows.

$$V^\pi(s) = R^\pi(s) + \gamma \sum_{\forall s'} \mathbf{P}_{ss'}^\pi V^\pi(s') \quad (\text{E1, p18})$$

- It means that, given a  $\pi$ , its value is determined by immediate reward plus discounted sum of future reward.
- In this light, let's try to criticize the  $\pi_{\text{speed}}$ .

```
t(policy_eval(pi_speed))
```

```
##           0           10           20           30           40           50           60 70
## [1,] -5.805929 -5.208781 -4.139262 -3.475765 -2.35376 -1.735376 -1.673538 0
```

- From the state 60, current policy gives the estimate for the state-value function of -1.6735376.
- We know that switching to *normal mode* at state 60 is better alternative than current action of *speed mode*. Because it guarantess the arrival to the state 70 with additional energy spending of 1.0.
- How would you express this fact in a mathematical form?



- Under the current policy's ( $\pi_{speed}$ ) value function, on the state 60,
  - Choosing *normal mode* gives

$$R + \gamma \mathbf{P}V = -1.0 + 1.0 \cdot 0 = -1.0$$

- Choosing *speed mode* gives

$$R + \gamma \mathbf{P}V = -1.5 + (0.9 \cdot 0 + 0.1 \cdot -2.6) = -1.76$$

- This,  $\pi_{speed}$  should modify its action on the state 60.
- You just improved the current policy  $\pi_{speed}$  for the state 60!
- This should be checked for all states as well as the state 60.
- This completes **policy improvement**.
- Formally, policy improvement implies the following task of replacement:

$$\pi^{new}(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} \left[ R(s, a) + \gamma \sum_{\forall s'} P_{ss'}^a V^{\pi^{old}}(s') \right], \text{ for all } s$$

$$\pi^{new}(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} \left[ R(s, a) + \gamma \sum_{\forall s'} P_{ss'}^a V^{\pi^{old}}(s') \right], \text{ for all } s$$

- The term in the RHS,  $R(s, a) + \gamma \sum_{\forall s'} P_{ss'}^a V^{\pi^{old}}(s')$ , implies **[an expected return of starting from state  $s$ , choosing an action  $a$  for this time step only, then following the policy  $\pi$  afterwards.]**
- How is this quantity different from  $V^{\pi}(s)$ ?
- The RHS makes an improvement using current policy  $\pi$ , by varying only the action in this time step.
- Formally,  $q(s, a)$  is called **action-value function**, also famously known as *Q-function*.

$$\begin{aligned} q(s, a) &:= \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a] \\ &= R(s, a) + \gamma \sum_{\forall s'} P_{ss'}^a V^{\pi^{old}}(s') \end{aligned}$$

$$\pi^{new}(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} \left[ R(s, a) + \gamma \sum_{\forall s'} P_{ss'}^a V^{\pi^{old}}(s') \right], \text{ for all } s$$

- Using this new notation of  $q(s, a)$ , the policy improvement can be written as

$$\pi^{new}(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} q(s, a), \text{ for all } s$$

- The improvement is called *greedy improvement* since it involves a myopic digression from the current policy, in a way that an action only on this time step is revised.
- It can be proved that *greedy improvement* is guaranteed to improve.

# Implementation

$$\pi^{new}(s) \leftarrow \operatorname{argmax}_{a \in \mathcal{A}} \left[ R(s, a) + \gamma \sum_{\forall s'} \mathbf{P}_{ss'}^a V^{\pi^{old}}(s') \right], \text{ for all } s$$

```
V_old <- policy_eval(pi_speed)
pi_old <- pi_speed
q_s_a <- R_s_a +
  cbind(gamma*P_normal%%V_old,
        gamma*P_speed%%V_old)
q_s_a
```

```
##      normal      speed
## 0 -6.208781 -5.805929
## 10 -5.139262 -5.208781
## 20 -4.475765 -4.139262
## 30 -3.353760 -3.475765
## 40 -1.735376 -2.353760
## 50 -2.673538 -1.735376
## 60 -1.000000 -1.673538
## 70 0.000000 0.000000
```

```
pi_new_vec <- apply(q_s_a, 1, which.max)
pi_new <- array(0, dim = dim(pi_old),
               dimnames = dimnames(pi_old))
for (i in 1:length(pi_new_vec)) {
  pi_new[i, pi_new_vec[i]] <- 1
}
pi_new
```

```
##      normal speed
## 0      0      1
## 10     1      0
## 20     0      1
## 30     1      0
## 40     1      0
## 50     0      1
## 60     1      0
## 70     1      0
```

● `policy_improve()`

```

policy_improve <- function(
  V_old,
  pi_old = pi_old, R_s_a = R_s_a, gamma = gamma,
  P_normal = P_normal, P_speed = P_speed) {

  q_s_a <- R_s_a + cbind(gamma*P_normal**V_old,
                        gamma*P_speed**V_old)

  pi_new_vec <- apply(q_s_a, 1, which.max)
  pi_new <- array(0, dim = dim(pi_old),
                 dimnames = dimnames(pi_old))

  for (i in 1:length(pi_new_vec)) {
    pi_new[i, pi_new_vec[i]] <- 1
  }
  return(pi_new)
}

```

● One step improvement from  $\pi^{speed}$ 

```

pi_old <- pi_speed
V_old <- policy_eval(pi_old)
pi_new <- policy_improve(V_old,
  pi_old = pi_old, R_s_a = R_s_a, gamma = gamma,
  P_normal = P_normal, P_speed = P_speed)

```

pi\_old

```

##      normal speed
## 0      0      1
## 10     0      1
## 20     0      1
## 30     0      1
## 40     0      1
## 50     0      1
## 60     0      1
## 70     0      1

```

pi\_new

```

##      normal speed
## 0      0      1
## 10     1      0
## 20     0      1
## 30     1      0
## 40     1      0
## 50     0      1
## 60     1      0
## 70     1      0

```

## III. Policy iteration

## Discussion

- Given a policy  $\pi$ , `policy_eval()` evaluates its state-value function.
- Using the estimate of state-value function, `policy_improe()` improves the policy to the better one.
- If this process is iterated, then it is guaranteed to reach optimal policy.
- In other words, policy iteration is the process to reach the optimal policy described as follows.

$$\begin{array}{ccccccc} \pi_0 & \xrightarrow{\text{policy\_eval()}} & V_0 & \xrightarrow{\text{policy\_improve()}} & \pi_1 & \xrightarrow{\text{policy\_eval()}} & V_1 \xrightarrow{\text{policy\_improve()}} \\ \pi_2 & \rightarrow \dots \rightarrow \dots \rightarrow & \pi^* & \xrightarrow{\text{policy\_eval()}} & V^*(= V^{\pi^*}, \text{ for short.}) \end{array}$$

- The iteration process terminates when  $\pi_i$  does not change any more, i.e.  $\pi_i = \pi_{i+1}$ .
- Note that policy evaluation is an approximate algorithm. For policy iteration purpose, policy evaluation cannot be, (and doesn't have to be as well), perfect.

# Try do it over and over until no change - from $\pi^{speed}$

## ● Step 0

```
pi_old <- pi_speed  
pi_old
```

```
##      normal speed  
## 0         0      1  
## 10        0      1  
## 20        0      1  
## 30        0      1  
## 40        0      1  
## 50        0      1  
## 60        0      1  
## 70        0      1
```

## ● Step 1

```
V_old <- policy_eval(pi_old)  
pi_new <- policy_improve(V_old,  
  pi_old = pi_old, R_s_a = R_s_a, gamma = gamma,  
  P_normal = P_normal, P_speed = P_speed)  
pi_old <- pi_new  
pi_old
```

```
##      normal speed  
## 0         0      1  
## 10        1      0  
## 20        0      1  
## 30        1      0  
## 40        1      0  
## 50        0      1  
## 60        1      0  
## 70        1      0
```



## ● Step 2

```
V_old <- policy_eval(pi_old)
pi_new <- policy_improve(V_old,
  pi_old = pi_old, R_s_a = R_s_a, gamma = gamma,
  P_normal = P_normal, P_speed = P_speed)
pi_old <- pi_new
pi_old
```

##	normal	speed
## 0	0	1
## 10	0	1
## 20	0	1
## 30	1	0
## 40	1	0
## 50	0	1
## 60	1	0
## 70	1	0

## ● Step 3

```
V_old <- policy_eval(pi_old)
pi_new <- policy_improve(V_old,
  pi_old = pi_old, R_s_a = R_s_a, gamma = gamma,
  P_normal = P_normal, P_speed = P_speed)
pi_old <- pi_new
pi_old
```

##	normal	speed
## 0	0	1
## 10	0	1
## 20	0	1
## 30	1	0
## 40	1	0
## 50	0	1
## 60	1	0
## 70	1	0

## Policy iteration process (from $\pi^{\text{speed}}$ )

- Now we are ready to implement whole process as a single code block.

```
pi_old <- pi_speed
cnt <- 0
repeat{ # do-while in R
  print(paste0(cnt, "-th iteration"))
  print(t(pi_old))
  V_old <- policy_eval(pi_old)
  pi_new <- policy_improve(V_old,
    pi_old = pi_old, R_s_a = R_s_a, gamma = gamma,
    P_normal = P_normal, P_speed = P_speed)
  if (all.equal(pi_new, pi_old)==TRUE) break
  pi_old <- pi_new
  cnt <- cnt + 1
}
print(policy_eval(pi_new))
```

```
## [1] "0-th iteration"
##           0 10 20 30 40 50 60 70
## normal 0  0  0  0  0  0  0  0
## speed  1  1  1  1  1  1  1  1
## [1] "1-th iteration"
##           0 10 20 30 40 50 60 70
## normal 0  1  0  1  1  0  1  1
## speed  1  0  1  0  0  1  0  0
## [1] "2-th iteration"
##           0 10 20 30 40 50 60 70
## normal 0  0  0  1  1  0  1  1
## speed  1  1  1  0  0  1  0  0
##           [,1]
## 0  -5.107744
## 10 -4.410774
## 20 -3.441077
## 30 -2.666667
## 40 -1.666667
## 50 -1.666667
## 60 -1.000000
## 70  0.000000
```

## Policy iteration process (from $\pi^{50}$ )

- The process should work for other initial choice of  $\pi$ , albeit possibly different convergence rate.

```
pi_old <- pi_50
cnt <- 0
repeat{ # do-while in R
  print(paste0(cnt, "-th iteration"))
  print(t(pi_old))
  V_old <- policy_eval(pi_old)
  pi_new <- policy_improve(V_old,
    pi_old = pi_old, R_s_a = R_s_a, gamma = gamma,
    P_normal = P_normal, P_speed = P_speed)
  if (all.equal(pi_new, pi_old)==TRUE) break
  pi_old <- pi_new
  cnt <- cnt + 1
}
print(policy_eval(pi_new))
```

```
## [1] "0-th iteration"
##           0  10  20  30  40  50  60  70
## normal 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5
## speed  0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5
## [1] "1-th iteration"
##           0  10  20  30  40  50  60  70
## normal 0   1   0   1   1   0   1   1
## speed  1   0   1   0   0   1   0   0
## [1] "2-th iteration"
##           0  10  20  30  40  50  60  70
## normal 0   0   0   1   1   0   1   1
## speed  1   1   1   0   0   1   0   0
##           [,1]
## 0  -5.107744
## 10 -4.410774
## 20 -3.441077
## 30 -2.666667
## 40 -1.666667
## 50 -1.666667
## 60 -1.000000
## 70  0.000000
```

"Success isn't permanent, and failure isn't fatal. - Mike Ditka"