

Lecture F4. MDP without Model 4

Sim, Min Kyu, Ph.D., mksim@seoultech.ac.kr



서울과학기술대학교 데이터사이언스학과

- 1 I. Policy iteration 3 - Q-learning control
- 2 II. Policy iteration 4 - Double Q-learning control

- `skier.R` is loaded as follows.

```
source("../skier.R")
```

```
## [1] "Skier's problem is set."  
## [1] "Defined are `state`, `P_normal`, `P_speed`, `R_s_a`, `q_s_a_init` (F2, p15)."  
## [1] "Defined are `pi_speed`, and `pi_50` (F2, p16)."  
## [1] "Defined are `simul_path`() (F2, p17)."  
## [1] "Defined are `simul_step`() (F2, p18)."  
## [1] "Defined are `pol_eval_MC`() (F2, p19)."  
## [1] "Defined are `pol_eval_TD`() (F2, p20)."  
## [1] "Defined are `pol_imp`() (F2, p20)."
```

I. Policy iteration 3 - Q-learning control

Introduction

- (`pol_eval_MC()`) MC control updates $q(s, a)$:

$$q(s, a) \leftarrow q(s, a) + \alpha(G_t - q(s, a)), \quad \forall s, a$$

- (`pol_eval_TD()`) TD control updates $q(s, a)$:

$$q(s, a) \leftarrow q(s, a) + \alpha(r_t + \gamma q(s', a') - q(s, a)), \quad \forall s, a$$

- (`pol_eval_Q()`) Q-learning updates $q(s, a)$:

$$q(s, a) \leftarrow q(s, a) + \alpha(r_t + \gamma \max_{a' \in \mathcal{A}} q(s', a') - q(s, a)), \quad \forall s, a$$

- Q-learning is a variation of TD control.
- Q-learning is greedy in a sense by taking maximum among possible future action.
- It is called **off-policy learning** since it may take action other than the current policy dictates.

Write `pol_eval_Q()`

```
pol_eval_TD <- function(sample_step, q_s_a, alpha) {
  s <- sample_step[1]
  a <- sample_step[2]
  r <- sample_step[3] %>% as.numeric()
  s_next <- sample_step[4]
  a_next <- sample_step[5]
  q_s_a[s,a] <- q_s_a[s,a] + alpha*(r+q_s_a[s_next,a_next]-q_s_a[s,a])
  return(q_s_a)
}
```

```
pol_eval_Q <- function(sample_step, q_s_a, alpha) {
  s <- sample_step[1]
  a <- sample_step[2]
  r <- sample_step[3] %>% as.numeric()
  s_next <- sample_step[4]
  a_next <- sample_step[5] # not used here
  q_s_a[s,a] <- q_s_a[s,a] + alpha*(r+max(q_s_a[s_next,])-q_s_a[s,a]) # change here
  return(q_s_a)
}
```

Q-learning

```

num_ep <- 10^5
beg_time <- Sys.time()
q_s_a <- q_s_a_init
pi <- pi_50
exploration_rate <- 1
for (epi_i in 1:num_ep) {
  s_now <- "0"
  while (s_now != "70") {
    sample_step <- simul_step(pi, s_now, P_normal, P_speed, R_s_a)
    q_s_a <- pol_eval_Q(sample_step, q_s_a, alpha = max(1/epi_i, 0.05))
    if (epi_i %% 100 == 0) {
      pi <- pol_imp(pi, q_s_a, epsilon = exploration_rate)
    }
    s_now <- sample_step[4]
    exploration_rate <- max(exploration_rate*0.9995, 0.001)
  }
}
end_time <- Sys.time()
t(q_s_a)

```

```
print(end_time-beg_time)
```

```
## Time difference of 47.5 secs
```

```
t(pi)
```

```
##    0 10 20 30 40 50 60 70
```

```
## n 0  0  1  1  0  1  1  1
```

```
## s 1  1  0  0  1  0  0  0
```

```

##          0          10          20          30          40          50          60 70
## n -5.713 -4.760 -3.746 -2.736 -1.999 -2.000 -1.000  0
## s -5.668 -4.701 -3.947 -3.237 -1.677 -2.027 -1.815  0

```

II. Policy iteration 4 - Double Q-learning control

Method

- (`pol_eval_Q()`) Q-learning updates $q(s, a)$:

$$q(s, a) \leftarrow q(s, a) + \alpha(r_t + \gamma \max_{a' \in \mathcal{A}} q(s', a') - q(s, a)), \quad \forall s, a$$

- (`pol_eval_dbl_Q()`) Double Q-learning uses a pair of q-functions, $q_1()$ and $q_2()$. It updates

- with probability 0.5

$$q_1(s, a) \leftarrow q_1(s, a) + \alpha(r_t + \gamma q_2(s', \operatorname{argmax}_{a' \in \mathcal{A}} q_1(s', a')) - q_1(s, a)), \quad \forall s, a$$

- with probability 0.5

$$q_2(s, a) \leftarrow q_2(s, a) + \alpha(r_t + \gamma q_1(s', \operatorname{argmax}_{a' \in \mathcal{A}} q_2(s', a')) - q_2(s, a)), \quad \forall s, a$$

- Policy is improved using $q_1(\cdot, \cdot) + q_2(\cdot, \cdot)$.

Write `pol_eval_dbl_Q()`

```
pol_eval_Q <- function(sample_step, q_s_a, alpha) {
  s <- sample_step[1]
  a <- sample_step[2]
  r <- sample_step[3] %>% as.numeric()
  s_next <- sample_step[4]
  q_s_a[s,a] <- q_s_a[s,a] + alpha*(r+max(q_s_a[s_next,])-q_s_a[s,a]) # change here
  return(q_s_a)
}

pol_eval_dbl_Q <- function(sample_step, q_s_a_1, q_s_a_2, alpha) {
  s <- sample_step[1]
  a <- sample_step[2]
  r <- sample_step[3] %>% as.numeric()
  s_next <- sample_step[4]
  if (runif(1) < 0.5) { # update q_s_a_1
    q_s_a_1[s,a] <- q_s_a_1[s,a] +
      alpha*(r+q_s_a_2[s_next, which.max(q_s_a_1[s_next,])]-q_s_a_1[s,a]) # change here
  } else { # update q_s_a_2
    q_s_a_2[s,a] <- q_s_a_2[s,a] +
      alpha*(r+q_s_a_1[s_next, which.max(q_s_a_2[s_next,])]-q_s_a_2[s,a]) # change here
  }
  return(list(q_s_a_1, q_s_a_2))
}
```

Double Q-learning

```

num_ep <- 10^5
beg_time <- Sys.time() # change below
q_s_a_1 <- q_s_a_init; q_s_a_2 <- q_s_a_init
pi <- pi_50
exploration_rate <- 1
for (epi_i in 1:num_ep) {
  s_now <- "0"
  while (s_now != "70") {
    sample_step <- simul_step(pi, s_now, P_normal, P_speed, R_s_a)
    q_s_a <- pol_eval_dbl_Q(sample_step, q_s_a_1, q_s_a_2, alpha = max(1/epi_i, 0.05)) # change here
    q_s_a_1 <- q_s_a[[1]]; q_s_a_2 <- q_s_a[[2]] # change here
    if (epi_i % 100 == 0) {
      pi <- pol_imp(pi, q_s_a_1+q_s_a_2, epsilon = exploration_rate) # change here
    }
    s_now <- sample_step[4]
    exploration_rate <- max(exploration_rate*0.9995, 0.001)
  }
}
t((q_s_a_1 + q_s_a_2)/2)

##      0      10      20      30      40      50      60 70
## n -5.69 -4.884 -3.677 -2.665 -1.650 -1.982 -1.000 0
## s -5.28 -4.639 -3.886 -3.303 -1.807 -1.617 -1.594 0

```

```

print(end_time-beg_time)

## Time difference of 51.71 secs

t(pi)

##      0 10 20 30 40 50 60 70
## n 0  0  1  1  1  0  1  1
## s 1  1  0  0  0  1  0  0

```

Exercise 1

Feel free to try different schemes for the number of iterations and exploration decaying scenarios.

"It's not that I'm so smart, it's just that I stay with problems longer. - A. Einstein"