

Lecture E1. MDP with Model 1

Baek, Jong min

2021-01-29

차 례

Policy evaluation 1(Page21)	2
rewritten with intermediate saving(Page22)	3
visualization	5
policy evaluation 2	8
Test 2	11

Policy evaluation 1(Page21)

```
import numpy as np
import pandas as pd
R = np.hstack((np.repeat(-1.5, 4), -0.5, np.repeat(-1.5, 2), 0)).reshape(-1,1)
states = np.arange(0, 80, 10)
p = np.matrix([
    [0.1,0,0.9,0,0,0,0,0],
    [0.1,0,0,0.9,0,0,0,0],
    [0,0.1,0,0,0.9,0,0,0],
    [0,0,0.1,0,0,0.9,0,0],
    [0,0,0,0.1,0,0,0.9,0],
    [0,0,0,0,0.1,0,0,0.9],
    [0,0,0,0,0,0.1,0,0.9],
    [0,0,0,0,0,0,0,1.0]
])
p = pd.DataFrame(p,index=states,columns=states)
print(R)
```

```
## [-1.5]
## [-1.5]
## [-1.5]
## [-1.5]
## [-0.5]
## [-1.5]
## [-1.5]
## [ 0. ]
```

```
print(p)
```

```
##      0      10      20      30      40      50      60      70
## 0    0.1    0.0    0.9    0.0    0.0    0.0    0.0    0.0
## 10   0.1    0.0    0.0    0.9    0.0    0.0    0.0    0.0
## 20   0.0    0.1    0.0    0.0    0.9    0.0    0.0    0.0
## 30   0.0    0.0    0.1    0.0    0.0    0.9    0.0    0.0
## 40   0.0    0.0    0.0    0.1    0.0    0.0    0.9    0.0
## 50   0.0    0.0    0.0    0.0    0.1    0.0    0.0    0.9
## 60   0.0    0.0    0.0    0.0    0.0    0.1    0.0    0.9
## 70   0.0    0.0    0.0    0.0    0.0    0.0    0.0    1.0
```

rewritten with intermediate saving(Page22)

```
gamma=1.0
epsilon=10**-8
v_old = np.zeros(8).reshape(8,1)
v_new = R + np.dot(gamma*p,v_old)
while np.max(np.abs(v_new-v_old)) > epsilon :
    v_old = v_new
    v_new = R + np.dot(gamma*p,v_old)
print(v_new.T)

## [[-5.80592905 -5.2087811 -4.13926239 -3.47576467 -2.35376031 -1.73537603
##    -1.6735376  0.          ]]
```

```
gamma = 1.0
epsilon = 10**-8
v_old = np.zeros(8).reshape(8,1)
v_new = R + np.dot(gamma*p,v_old)
results = v_old.T
results = np.vstack([results,v_new.T])
while np.max(np.abs(v_new-v_old)) > epsilon :
    v_old = v_new
    v_new = R + np.dot(gamma*p,v_old)
    results = np.vstack([results,v_new.T])
print(v_new.T)

## [[-5.80592905 -5.2087811 -4.13926239 -3.47576467 -2.35376031 -1.73537603
##    -1.6735376  0.          ]]
```

```
results = pd.DataFrame(results,columns=states)
print(results.head())
```

```
##      0      10      20      30      40      50      60      70
## 0  0.000  0.0000  0.0000  0.000  0.000  0.0000  0.000  0.0
## 1 -1.500 -1.5000 -1.5000 -1.500 -0.500 -1.5000 -1.500  0.0
## 2 -3.000 -3.0000 -2.1000 -3.000 -2.000 -1.5500 -1.650  0.0
## 3 -3.690 -4.5000 -3.6000 -3.105 -2.285 -1.7000 -1.655  0.0
## 4 -5.109 -4.6635 -4.0065 -3.390 -2.300 -1.7285 -1.670  0.0
```

```
print(results.tail())
```

```
##      0      10      20      30      40      50      60      70
```

```
## 18 -5.805929 -5.208781 -4.139262 -3.475765 -2.35376 -1.735376 -1.673538 0.0
## 19 -5.805929 -5.208781 -4.139262 -3.475765 -2.35376 -1.735376 -1.673538 0.0
## 20 -5.805929 -5.208781 -4.139262 -3.475765 -2.35376 -1.735376 -1.673538 0.0
## 21 -5.805929 -5.208781 -4.139262 -3.475765 -2.35376 -1.735376 -1.673538 0.0
## 22 -5.805929 -5.208781 -4.139262 -3.475765 -2.35376 -1.735376 -1.673538 0.0
```


visualization

```
fig1=results[results.index < 6]
fig2=results[(results.index >= 7)&(results.index < 12)]
fig3=results[(results.index >= 13)&(results.index < 18)]
```

```
plt.plot(fig1.T,marker='o')
```

```
## [<matplotlib.lines.Line2D object at 0x000000002CB16278>, <matplotlib.lines.Line2D object at 0x000000002CB163>]
```


```
plt.legend(fig1.columns,title='states')
plt.grid(True)
plt.xlabel('state')
plt.ylabel('value_fn')
plt.title('Iteration from 1 to 6')
plt.show()
```



```
plt.plot(fig2.T,marker='o')
```

```
## [<matplotlib.lines.Line2D object at 0x000000002DC04F60>, <matplotlib.lines.Line2D object at 0x000000002DC04F60>]
```


```
plt.legend(fig1.columns,title='states')
plt.grid(True)
plt.xlabel('state')
plt.ylabel('value_fn')
plt.title('Iteration from 7 to 12')
plt.show()
```



```
plt.plot(fig3.T,marker='o')
```

```
## [<matplotlib.lines.Line2D object at 0x000000002DBED978>, <matplotlib.lines.Line2D object at 0x000000002DBED978>]
```

```
plt.legend(fig1.columns,title='states')
plt.grid(True)
plt.xlabel('state')
plt.ylabel('value_fn')
plt.title('Iteration from 13 to 18')
plt.show()
```



policy evaluation 2

$$\pi : S \rightarrow A$$

```
states = np.arange(0,80,10)
pi_speed = np.array([np.repeat(0,len(states)),np.repeat(1,len(states))]).T
pi_speed = pd.DataFrame(pi_speed,columns=['normal','speed'],index=[states])
pi_speed
```

```
##      normal  speed
## 0         0      1
## 10        0      1
## 20        0      1
## 30        0      1
## 40        0      1
## 50        0      1
## 60        0      1
## 70        0      1
```

$$R^\pi : S \rightarrow \mathbb{R}$$

```
R_s_a = np.array([[ -1, -1, -1, -1, 0.0, -1, -1, 0], [-1.5, -1.5, -1.5, -1.5, -0.5, -1.5, -1.5, 0]]).T
R_s_a = pd.DataFrame(R_s_a,columns=['normal','speed'],index=[states])
R_s_a
```

```
##      normal  speed
## 0      -1.0  -1.5
## 10     -1.0  -1.5
## 20     -1.0  -1.5
## 30     -1.0  -1.5
## 40      0.0  -0.5
## 50     -1.0  -1.5
## 60     -1.0  -1.5
## 70      0.0   0.0
```



```
def reward_fn(given_pi):
    R_s_a = np.array([[ -1, -1, -1, -1, 0.0, -1, -1, 0], [-1.5, -1.5, -1.5, -1.5, -0.5, -1.5, -1.5, 0]]).T
    R_pi = np.sum(given_pi*R_s_a,axis=1)
    return R_pi
reward_fn(pi_speed)
```

```
## 0      -1.5
## 10     -1.5
## 20     -1.5
## 30     -1.5
## 40     -0.5
## 50     -1.5
## 60     -1.5
## 70      0.0
## dtype: float64
```

$$P^\pi : S \times A \rightarrow S$$

```
states = np.arange(0,80,10)
p_normal = pd.DataFrame(np.array([
0,1,0,0,0,0,0,0,
0,0,1,0,0,0,0,0,
0,0,0,1,0,0,0,0,
0,0,0,0,1,0,0,0,
0,0,0,0,0,1,0,0,
0,0,0,0,0,0,1,0,
0,0,0,0,0,0,0,1,
0,0,0,0,0,0,0,1,
0,0,0,0,0,0,0,1
]).reshape(8,8),index=states, columns=states)
p_speed = pd.DataFrame(np.array([
.1,0,.9,0,0,0,0,0,
.1,0,0,.9,0,0,0,0,
0,.1,0,0,.9,0,0,0,
0,0,.1,0,0,.9,0,0,
0,0,0,.1,0,0,.9,0,
0,0,0,0,.1,0,0,.9,
0,0,0,0,0,.1,0,.9,
0,0,0,0,0,0,0,1,
]).reshape(8,8),index=states, columns=states)
```

```
def transition(given_pi,states,p_normal,p_speed):
    p_out = pd.DataFrame(np.zeros(shape=(len(states),len(states))),index=states, columns=states)
    for s in range(len(states)) :
        action_dist = given_pi.iloc[s]
        p = action_dist['normal']*p_normal + action_dist['speed']*p_speed
        p_out.iloc[s] = p.iloc[s]
    return p_out
```

```
transition(pi_speed,states,p_normal,p_speed)
```

```
##      0      10      20      30      40      50      60      70
## 0    0.1    0.0    0.9    0.0    0.0    0.0    0.0    0.0
## 10   0.1    0.0    0.0    0.9    0.0    0.0    0.0    0.0
## 20   0.0    0.1    0.0    0.0    0.9    0.0    0.0    0.0
## 30   0.0    0.0    0.1    0.0    0.0    0.9    0.0    0.0
## 40   0.0    0.0    0.0    0.1    0.0    0.0    0.9    0.0
## 50   0.0    0.0    0.0    0.0    0.1    0.0    0.0    0.9
## 60   0.0    0.0    0.0    0.0    0.0    0.1    0.0    0.9
## 70   0.0    0.0    0.0    0.0    0.0    0.0    0.0    1.0
```

Test 2

```
pi_50 = pi_speed = pd.DataFrame(np.array([np.repeat(0.5,len(states)),np.repeat(0.5,len(states))]).T,columns=[  
pi_50
```

```
##      normal  speed  
## 0         0.5    0.5  
## 10        0.5    0.5  
## 20        0.5    0.5  
## 30        0.5    0.5  
## 40        0.5    0.5  
## 50        0.5    0.5  
## 60        0.5    0.5  
## 70        0.5    0.5
```

```
transition(pi_50,states,p_normal,p_speed)
```

```
##      0      10      20      30      40      50      60      70  
## 0  0.05  0.50  0.45  0.00  0.00  0.00  0.00  0.00  
## 10 0.05  0.00  0.50  0.45  0.00  0.00  0.00  0.00  
## 20 0.00  0.05  0.00  0.50  0.45  0.00  0.00  0.00  
## 30 0.00  0.00  0.05  0.00  0.50  0.45  0.00  0.00  
## 40 0.00  0.00  0.00  0.05  0.00  0.50  0.45  0.00  
## 50 0.00  0.00  0.00  0.00  0.05  0.00  0.50  0.45  
## 60 0.00  0.00  0.00  0.00  0.00  0.05  0.00  0.95  
## 70 0.00  0.00  0.00  0.00  0.00  0.00  0.00  1.00
```

E1.Rmd

```
"Hello"
```

```
## [1] "Hello"
```