

F1_Exercises

Kwon do yun

2021-01-25

차 례

Preparation (P. 9)	2
Simulator - π^{speed} (P. 11)	3
Simulator - π^{50} (P. 13)	4
Implementation 1 - π^{speed} (vectorized) (P. 17)	5
Implementation 2 - π^{speed} (running estimate) (P. 19)	6
Implementation 3 - π^{50} (vectorized) (P. 21)	7
Implementation 4 - π^{50} (running estimate) (P. 23)	8

Preparation (P. 9)

```
import numpy as np
import pandas as pd
states=np.arange(0,80,10).astype('str')
P_normal=pd.DataFrame(np.matrix([[0,1,0,0,0,0,0,0],
                                [0,0,1,0,0,0,0,0],
                                [0,0,0,1,0,0,0,0],
                                [0,0,0,0,1,0,0,0],
                                [0,0,0,0,0,1,0,0],
                                [0,0,0,0,0,0,1,0],
                                [0,0,0,0,0,0,0,1],
                                [0,0,0,0,0,0,0,1]]), index=states,columns=states)
P_speed=pd.DataFrame(np.matrix([[.1,0,.9,0,0,0,0,0],
                                [.1,0,.9,0,0,0,0,0],
                                [0,.1,0,.9,0,0,0,0],
                                [0,0,.1,0,.9,0,0,0],
                                [0,0,0,.1,0,.9,0,0],
                                [0,0,0,0,.1,0,.9,0],
                                [0,0,0,0,0,.1,0,.9],
                                [0,0,0,0,0,0,.1,0,1]]), index=states, columns=states)

R_s_a=pd.DataFrame(np.matrix([-1,-1,-1,-1,0.0,-1,-1,0,-1.5,-1.5,-1.5,-1.5,-0.5,-1.5,-1.5,0])).reshape(len(states),len(states))
R_s_a.T
```

```
##      0  10  20  30  40  50  60  70
## n -1.0 -1.0 -1.0 -1.0  0.0 -1.0 -1.0  0.0
## s -1.5 -1.5 -1.5 -1.5 -0.5 -1.5 -1.5  0.0
```

```
pi_speed=pd.DataFrame(np.c_[np.repeat(0,len(states)),np.repeat(1,len(states))], index=states, columns=['n','s'])
pi_50=pd.DataFrame(np.c_[np.repeat(0.5,len(states)), np.repeat(0.5,len(states))],index=states, columns=['n','s'])
pi_speed.T
```

```
##      0  10  20  30  40  50  60  70
## n  0   0   0   0   0   0   0   0
## s  1   1   1   1   1   1   1   1
```

```
pi_50.T
```

```
##      0  10  20  30  40  50  60  70
## n  0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5
## s  0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5
```

Simulator - π^{speed} (P. 11)

```

pi=pi_speed
np.random.seed(1234)
history=[]
MC_N=10000
for MC_i in range(MC_N):
    s_now='0'
    history_i=list(s_now)

    while s_now != '70' :
        if np.random.uniform(0,1) < pi.loc[s_now]['n']:
            a_now='n'
            P=P_normal
        else:
            a_now='s'
            P=P_speed

        r_now=str(R_s_a.loc[s_now][a_now])
        s_next=states[np.argmax(P.loc[s_now].cumsum() < np.random.uniform(0,1))].item()
        history_i.extend([a_now,r_now,s_next])
        s_now=s_next

    history.append(history_i)

history_speed=history
func=np.vectorize(lambda x: ','.join(x))
pd.Series(func(history_speed[:20]))

```

```

## 0      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 1      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 2      0,s,-1.5,0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1...
## 3      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 4      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 5      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 6      0,s,-1.5,20,s,-1.5,10,s,-1.5,30,s,-1.5,50,s,-1...
## 7      0,s,-1.5,20,s,-1.5,40,s,-0.5,30,s,-1.5,50,s,-1...
## 8      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 9      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 10     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 11     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 12     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70

```

```

## 13          0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 14          0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 15      0,s,-1.5,20,s,-1.5,10,s,-1.5,30,s,-1.5,50,s,-1...
## 16          0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 17          0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 18          0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 19          0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## dtype: object
##
## C:\Users\User\AppData\Local\R-MINI~1\envs\R-RETI~1\lib\site-packages\numpy\core\_asarray.py:83: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequence (which is a tuple or tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you should use np.concatenate.
##     return array(a, dtype, copy=False, order=order)

```

Simulator - π^{50} (P. 13)

```

pi=pi_50
np.random.seed(1234)
history=[]
MC_N=10000
for MC_i in range(MC_N):
    s_now='0'
    history_i=list(s_now)

    while s_now != '70' :
        if np.random.uniform(0,1) < pi.loc[s_now]['n']:
            a_now='n'
            P=P_normal
        else:
            a_now='s'
            P=P_speed

        r_now=str(R_s_a.loc[s_now][a_now])
        s_next=states[np.argmax(P.loc[s_now].cumsum() < np.random.uniform(0,1))].item()
        history_i.extend([a_now,r_now,s_next])
        s_now=s_next

    history.append(history_i)

history_50=history
func=np.vectorize(lambda x: ','.join(x))
pd.Series(func(history_50[:20]))

```

```

## 0      0,n,-1.0,10,n,-1.0,20,s,-1.5,40,n,0.0,50,s,-1....
## 1      0,n,-1.0,10,s,-1.5,30,n,-1.0,40,s,-0.5,30,s,-1...
## 2      0,s,-1.5,20,n,-1.0,30,n,-1.0,40,s,-0.5,60,s,-1...
## 3      0,s,-1.5,20,n,-1.0,30,n,-1.0,40,n,0.0,50,n,-1....
## 4      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,s,-1.5,20,s,-1...
## 5      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,n,-1.0,40,n,0....
## 6      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,n,-1.0,40,n,0....
## 7              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,n,-1.0,70
## 8      0,s,-1.5,20,n,-1.0,30,s,-1.5,50,n,-1.0,60,s,-1...
## 9      0,s,-1.5,20,s,-1.5,40,n,0.0,50,n,-1.0,60,n,-1....
## 10     0,n,-1.0,10,s,-1.5,30,n,-1.0,40,s,-0.5,60,s,-1...
## 11     0,s,-1.5,20,n,-1.0,30,n,-1.0,40,n,0.0,50,n,-1....
## 12     0,n,-1.0,10,s,-1.5,30,n,-1.0,40,n,0.0,50,s,-1....
## 13              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 14              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 15              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 16     0,s,-1.5,20,s,-1.5,40,n,0.0,50,n,-1.0,60,n,-1....
## 17              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,n,-1.0,70
## 18     0,n,-1.0,10,n,-1.0,20,s,-1.5,40,n,0.0,50,n,-1....
## 19              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## dtype: object
##
## C:\Users\User\AppData\Local\R-MINI~1\envs\R-RETI~1\lib\site-packages\numpy\core\_asarray.py:83: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequence (which is a tuple-or-list of tuple-or-lists or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you should use 'numpy.array' from the numpy.compat module instead.
## return array(a, dtype, copy=False, order=order)

```

Implementation 1 - π^{speed} (vectorized) (P. 17)

```

pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','sum'])
pol_eval.T

```

```

##           0   10   20   30   40   50   60   70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## sum    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0

```

```

for MC_i in range(len(history_speed)):
    history_i=history_speed[MC_i]

    for j in range(0,len(history_i),3):
        pol_eval.loc[history_i[j]]['count']+=1

```

```

        if j < len(history_i) :
            pol_eval.loc[history_i[j]]['sum']+=pd.Series(history_i)[range(j+2,len(history_i)-1,3)].astype('float64')

        else:
            pol_eval.loc[history_i[j]]['sum']+=0
pol_eval.T

```

```

##           0      10      20      30      40      50      60      70
## count  11225.0  1076.0  10291.0  1887.0   9485.0  2563.0   8563.0  10000.0
## sum    -65136.0 -5619.0 -42703.0 -6539.0 -22275.5 -4472.5 -14355.0    0.0

```

```

pol_eval['sum']/pol_eval['count']

```

```

## 0      -5.802762
## 10     -5.222119
## 20     -4.149548
## 30     -3.465289
## 40     -2.348498
## 50     -1.745025
## 60     -1.676398
## 70      0.000000
## dtype: float64

```

Implementation 2 - π^{speed} (running estimate) (P. 19)

```

pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','est'])
pol_eval.T

```

```

##           0   10   20   30   40   50   60   70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0

```

```

for MC_i in range(len(history_speed)):
    history_i=history_speed[MC_i]

    for j in range(0,len(history_i),3):
        # update count
        pol_eval.loc[history_i[j]]['count']+=1
        current_cnt=pol_eval.loc[history_i[j]]['count']

        # return is the new info

```

```

if j < len(history_i):
    new_info=pd.Series(history_i)[range(j+2,len(history_i)-1,3)].astype('float').sum()

else:
    new_info=0

# update the last estimate with new info
alpha=1/current_cnt
pol_eval.loc[history_i[j]]['est']+=alpha*(new_info-pol_eval.loc[history_i[j]]['est'])

np.round(pol_eval.T,2)

```

```

##           0         10         20         30         40         50         60         70
## count  11225.0  1076.00  10291.00  1887.00  9485.00  2563.00  8563.00  10000.0
## est    -5.8    -5.22    -4.15    -3.47    -2.35    -1.75    -1.68     0.0

```

Implementation 3 - π^{50} (vectorized) (P. 21)

```

pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','sum'])
pol_eval.T

```

```

##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## sum    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0

```

```

for MC_i in range(len(history_50)):
    history_i=history_50[MC_i]

    for j in range(0,len(history_i),3):
        pol_eval.loc[history_i[j]]['count']+=1

        if j < len(history_i) :
            pol_eval.loc[history_i[j]]['sum']+=pd.Series(history_i)[range(j+2,len(history_i)-1,3)].astype('float').sum()

        else:
            pol_eval.loc[history_i[j]]['sum']+=0

pol_eval.T

```

```

##           0         10         20         30         40         50         60         70
## count  10863.0   5792.0   8140.0   7121.0   7549.0   7363.0   6991.0  10000.0

```

```
## sum    -64904.5 -29662.5 -33549.0 -24133.0 -15410.0 -14874.5 -9436.5      0.0
```

```
pol_eval['sum']/pol_eval['count']
```

```
## 0      -5.974823
## 10     -5.121288
## 20     -4.121499
## 30     -3.388990
## 40     -2.041330
## 50     -2.020168
## 60     -1.349807
## 70      0.000000
## dtype: float64
```

Implementation 4 - π^{50} (running estimate) (P. 23)

```
pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','est'])
pol_eval.T
```

```
##          0   10   20   30   40   50   60   70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(len(history_50)):
    history_i=history_50[MC_i]

    for j in range(0,len(history_i),3):
        # increment count
        pol_eval.loc[history_i[j]]['count']+=1
        current_cnt=pol_eval.loc[history_i[j]]['count']

        # return is the new info
        if j < len(history_i):
            new_info=pd.Series(history_i)[range(j+2,len(history_i)-1,3)].astype('float').sum()

        else:
            new_info=0

        # update the last estimate with new info
        alpha=1/current_cnt
        pol_eval.loc[history_i[j]]['est']+=alpha*(new_info-pol_eval.loc[history_i[j]]['est'])
```



```
np.round(pol_eval.T,2)
```

##	0	10	20	30	40	50	60	70
## count	10863.00	5792.00	8140.00	7121.00	7549.00	7363.00	6991.00	10000.0
## est	-5.97	-5.12	-4.12	-3.39	-2.04	-2.02	-1.35	0.0

```
"F1_Exercises"
```