

F1 Python ver

Lee SungHo

2021-02-01

Contents

page 9 Preparation	2
page 11 simulator pi speed	3
page 13 simulator pi 50	5
page 17 Implementation 1 π^{speed} (vectorized)	7
page 19 Implementation 2 π^{speed} (running estimate)	8
page 21 Implementation 3 π^{50} (vectorized)	9
page 23 Implementation 4 π^{50} (running estimate)	10
page 36 Implementation 5 pi speed	11
page 37 Implementation 6 pi 50	12

page 9 Preparation

```
import numpy as np
import pandas as pd

states=np.arange(0,80,10).astype('str')
P_normal=pd.DataFrame(np.matrix([[0,1,0,0,0,0,0,0],
                                [0,0,1,0,0,0,0,0],
                                [0,0,0,1,0,0,0,0],
                                [0,0,0,0,1,0,0,0],
                                [0,0,0,0,0,1,0,0],
                                [0,0,0,0,0,0,1,0],
                                [0,0,0,0,0,0,0,1],
                                [0,0,0,0,0,0,0,1]]), index=states,columns=states)

P_speed=pd.DataFrame(np.matrix([[.1,0,.9,0,0,0,0,0],
                                [.1,0,0,.9,0,0,0,0],
                                [0,.1,0,0,.9,0,0,0],
                                [0,0,.1,0,0,.9,0,0],
                                [0,0,0,.1,0,0,.9,0],
                                [0,0,0,0,.1,0,0,.9],
                                [0,0,0,0,0,.1,0,.9],
                                [0,0,0,0,0,0,.1,1]]), index=states, columns=states)

R_s_a=pd.DataFrame(np.matrix([-1,-1,-1,-1,0.0,-1,-1,0,-1.5,-1.5,-1.5,-1.5,-0.5,-1.5,-1.5,0])).reshape(16,8)
print(R_s_a.T)

##      0   10   20   30   40   50   60   70
## n -1.0 -1.0 -1.0 -1.0  0.0 -1.0 -1.0  0.0
## s -1.5 -1.5 -1.5 -1.5 -0.5 -1.5 -1.5  0.0

pi_speed=pd.DataFrame(np.c_[np.repeat(0,len(states)),np.repeat(1,len(states))], index=states, columns=states)

pi_50=pd.DataFrame(np.c_[np.repeat(0.5,len(states)), np.repeat(0.5,len(states))],index=states, columns=states)

print(pi_speed.T)

##      0   10   20   30   40   50   60   70
## n  0    0    0    0    0    0    0    0
## s  1    1    1    1    1    1    1    1

print(pi_50.T)

##      0   10   20   30   40   50   60   70
## n  0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5
## s  0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5
```

page 11 simulator pi speed

```
pi=pi_speed
np.random.seed(1234)
history=[]
MC_N=10000

for MC_i in range(MC_N):
    s_now='0'
    history_i=list(s_now)

    while s_now != '70' :
        if np.random.uniform(0,1) < pi.loc[s_now]['n']:
            a_now='n'
            P=P_normal
        else:
            a_now='s'
            P=P_speed

        r_now=str(R_s_a.loc[s_now][a_now])
        s_next=states[np.argmin(np.cumsum(P.loc[s_now,])) < np.random.uniform(0,1))]
        history_i.extend([a_now,r_now,s_next])
        s_now=s_next

    history.append(history_i)

history_speed=history
func=np.vectorize(lambda x: ','.join(x))
pd.Series(func(history_speed[:20]))
```

```
## 0      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 1      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 2      0,s,-1.5,0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1...
## 3      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 4      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 5      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 6      0,s,-1.5,20,s,-1.5,10,s,-1.5,30,s,-1.5,50,s,-1...
## 7      0,s,-1.5,20,s,-1.5,40,s,-0.5,30,s,-1.5,50,s,-1...
## 8      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 9      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 10     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 11     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 12     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 13     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 14     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 15     0,s,-1.5,20,s,-1.5,10,s,-1.5,30,s,-1.5,50,s,-1...
## 16     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 17     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 18     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 19     0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## dtype: object
##
```

```
## C:\Users\LEESUN~1\AppData\Local\R-MINI~1\envs\R-RETI~1\lib\site-packages\numpy\core\_asarray.py:83: \
```

```
##    return array(a, dtype, copy=False, order=order)
```

page 13 simulator pi 50

```
pi=pi_50
np.random.seed(1234)
history=[]
MC_N=10000

for MC_i in range(MC_N):
    s_now='0'
    history_i=list(s_now)

    while s_now != '70' :
        if np.random.uniform(0,1) < pi.loc[s_now]['n']:
            a_now='n'
            P=P_normal
        else:
            a_now='s'
            P=P_speed

        r_now=str(R_s_a.loc[s_now][a_now])
        s_next=states[np.argmin(P.loc[s_now].cumsum() < np.random.uniform(0,1))].item()
        history_i.extend([a_now,r_now,s_next])
        s_now=s_next

    history.append(history_i)

history_50=history
func=np.vectorize(lambda x: ','.join(x))
pd.Series(func(history_50[:20]))
```

```
## 0      0,n,-1.0,10,n,-1.0,20,s,-1.5,40,n,0.0,50,s,-1...
## 1      0,n,-1.0,10,s,-1.5,30,n,-1.0,40,s,-0.5,30,s,-1...
## 2      0,s,-1.5,20,n,-1.0,30,n,-1.0,40,s,-0.5,60,s,-1...
## 3      0,s,-1.5,20,n,-1.0,30,n,-1.0,40,n,0.0,50,n,-1...
## 4      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,s,-1.5,20,s,-1...
## 5      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,n,-1.0,40,n,0...
## 6      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,n,-1.0,40,n,0...
## 7              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,n,-1.0,70
## 8      0,s,-1.5,20,n,-1.0,30,s,-1.5,50,n,-1.0,60,s,-1...
## 9      0,s,-1.5,20,s,-1.5,40,n,0.0,50,n,-1.0,60,n,-1...
## 10     0,n,-1.0,10,s,-1.5,30,n,-1.0,40,s,-0.5,60,s,-1...
## 11     0,s,-1.5,20,n,-1.0,30,n,-1.0,40,n,0.0,50,n,-1...
## 12     0,n,-1.0,10,s,-1.5,30,n,-1.0,40,n,0.0,50,s,-1...
## 13              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 14              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 15              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 16     0,s,-1.5,20,s,-1.5,40,n,0.0,50,n,-1.0,60,n,-1...
## 17              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,n,-1.0,70
## 18     0,n,-1.0,10,n,-1.0,20,s,-1.5,40,n,0.0,50,n,-1...
## 19              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## dtype: object
##
```

```
## C:\Users\LEESUN~1\AppData\Local\R-MINI~1\envs\R-RETI~1\lib\site-packages\numpy\core\_asarray.py:83: \
```

```
##    return array(a, dtype, copy=False, order=order)
```

page 17 Implementation 1 π^{speed} (vectorized)

```
pol_eval=pd.DataFrame(np.matrix(np.zeros((len(states)*2))).reshape(len(states),2), index=states, columns=states)
print(pol_eval.T)
```

```
##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## sum    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(MC_N):
    history_i=history_speed[MC_i]

    for j in range(0,len(history_i),3):
        pol_eval.loc[history_i[j]]['count']+=1

        if j < len(history_i) :
            pol_eval.loc[history_i[j]]['sum']+=pd.Series(history_i)[range(j+2,len(history_i)-1,3)].astype(float)

        else:
            pol_eval.loc[history_i[j]]['sum']+=0
```

```
print(pol_eval.T)
```

```
##           0      10      20      30      40      50      60      70
## count 11225.0 1076.0 10291.0 1887.0 9485.0 2563.0 8563.0 10000.0
## sum  -65136.0 -5619.0 -42703.0 -6539.0 -22275.5 -4472.5 -14355.0 0.0
```

```
pol_cal=pd.DataFrame(pol_eval['sum']/pol_eval['count'])
print(pol_cal.T)
```

```
##           0      10      20      30      40      50      60      70
## 0 -5.802762 -5.222119 -4.149548 -3.465289 -2.348498 -1.745025 -1.676398 0.0
```

page 19 Implementation 2 π^{speed} (running estimate)

```
pol_eval=pd.DataFrame(np.matrix(np.zeros((len(states)*2))).reshape(len(states),2), index=states, columns=states)
print(pol_eval.T)
```

```
##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(MC_N):
    history_i=history_speed[MC_i]

    for j in range(0,len(history_i),3):
        # update count
        pol_eval.loc[history_i[j]]['count']+=1
        current_cnt=pol_eval.loc[history_i[j]]['count']

        # return is the new info
        if j < len(history_i):
            new_info=pd.Series(history_i)[range(j+2,len(history_i)-1,3)].astype('float').sum()

        else:
            new_info=0

        # update the last estimate with new info
        alpha=1/current_cnt
        pol_eval.loc[history_i[j]]['est']+=alpha*(new_info-pol_eval.loc[history_i[j]]['est'])

print(pol_eval.T)
```

```
##           0           10    ...           60           70
## count  11225.000000  1076.000000  ...  8563.000000  10000.0
## est    -5.802762   -5.222119  ...   -1.676398    0.0
##
## [2 rows x 8 columns]
```


page 21 Implementation 3 π^{50} (vectorized)

```
pol_eval=pd.DataFrame(np.matrix(np.zeros((len(states)*2))).reshape(len(states),2), index=states, columns=states)
pol_eval.T
```

```
##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## sum    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(MC_N):
    history_i=history_50[MC_i]

    for j in range(0,len(history_i),3):
        pol_eval.loc[history_i[j]]['count']+=1

        if j < len(history_i) :
            pol_eval.loc[history_i[j]]['sum']+=pd.Series(history_i)[range(j+2,len(history_i)-1,3)].astype(float)

        else:
            pol_eval.loc[history_i[j]]['sum']+=0
pol_eval.T
```

```
##           0          10          20          30          40          50          60          70
## count 10863.0   5792.0   8140.0   7121.0   7549.0   7363.0   6991.0  10000.0
## sum  -64904.5 -29662.5 -33549.0 -24133.0 -15410.0 -14874.5 -9436.5      0.0
```

```
pol_cal=pd.DataFrame(pol_eval['sum']/pol_eval['count'])
print(pol_cal.T)
```

```
##           0          10          20          30          40          50          60          70
## 0 -5.974823 -5.121288 -4.121499 -3.38899 -2.04133 -2.020168 -1.349807  0.0
```

page 23 Implementation 4 π^{50} (running estimate)

```
pol_eval=pd.DataFrame(np.matrix(np.zeros((len(states)*2))).reshape(len(states),2), index=states, columns=states)
pol_eval.T
```

```
##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(MC_N):
    history_i=history_50[MC_i]

    for j in range(0,len(history_i),3):
        # increment count
        pol_eval.loc[history_i[j]]['count']+=1
        current_cnt=pol_eval.loc[history_i[j]]['count']

        # return is the new info
        if j < len(history_i):
            new_info=pd.Series(history_i)[range(j+2,len(history_i)-1,3)].astype('float').sum()

        else:
            new_info=0

        # update the last estimate with new info
        alpha=1/current_cnt
        pol_eval.loc[history_i[j]]['est']+=alpha*(new_info-pol_eval.loc[history_i[j]]['est'])

print(pol_cal.T)
```

```
##           0           10           20           30           40           50           60    70
## 0 -5.974823 -5.121288 -4.121499 -3.38899 -2.04133 -2.020168 -1.349807  0.0
```

page 36 Implementation 5 pi speed

```
import pandas as pd
import numpy as np

pol_eval=pd.DataFrame(np.matrix(np.zeros((len(states)*2))).reshape(len(states),2), index=states, columns=states)

print(pol_eval.T)

##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0

for episode_i in range(len(history_speed)):
    history_i = history_speed[episode_i]

    # update count
    for j in range(0,len(history_i),3):
        pol_eval.loc[history_i[j]]['count'] +=1
        current_cnt =pol_eval.loc[history_i[j]]['count']

    #build TD target
    if(j < len(history_i)-3):
        TD_tgt = float(history_i[j+2])+pol_eval.loc[history_i[j+3]]['est']

    else:
        TD_tgt = 0

    # TD-updating
    alpha = 1/current_cnt

    pol_eval.loc[history_i[j]]['est'] += alpha*(TD_tgt - pol_eval.loc[history_i[j]]['est'])

print(pol_eval.T)

##           0           10    ...           60           70
## count  11225.000000  1076.000000  ...  8563.000000  10000.0
## est    -5.738838    -5.186466  ...    -1.675699    0.0
##
## [2 rows x 8 columns]
```

page 37 Implementation 6 pi 50

```
pol_eval=pd.DataFrame(np.matrix(np.zeros((len(states)*2))).reshape(len(states),2), index=states, columns=states)
print(pol_eval.T)
```

```
##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for episode_i in range(len(history_50)):
    history_i = history_50[episode_i]

    # update count
    for j in range(0,len(history_i),3):
        pol_eval.loc[history_i[j]]['count'] +=1
        current_cnt = pol_eval.loc[history_i[j]]['count']

        #build TD target
        if(j < len(history_i)-3):
            TD_tgt = float(history_i[j+2])+pol_eval.loc[history_i[j+3]]['est']

        else:
            TD_tgt = 0

        # TD-updating
        alpha = 1/current_cnt

        pol_eval.loc[history_i[j]]['est'] += alpha*(TD_tgt - pol_eval.loc[history_i[j]]['est'])

print(pol_eval.T)
```

```
##           0           10           20    ...           50           60           70
## count  10863.00000  5792.000000  8140.000000  ...  7363.000000  6991.000000  10000.0
## est    -5.84492    -5.052485    -4.079273  ...    -2.026683    -1.351198    0.0
##
## [2 rows x 8 columns]
```