

F1 Solution

Reinforcement Learning Study

2021-02-21

차 례

Preparation	2
Simulator - π^{speed}	4
Simulator - π^{50}	6
Implementation 1 - π^{speed} (vectorized)	8
Implementation 2 - π^{speed} (running estimate)	9
Implementation 3 - π^{50} (vectorized)	10
Implementation 4 - π^{50} (running estimate)	11
TD - Implementation 5 - π^{speed}	12
TD - Implementation 6 - π^{50}	13

Preparation

```
states=np.arange(0,70+10,10).astype('str')
```

```
P_normal=pd.DataFrame(np.matrix([[0,1,0,0,0,0,0,0],
                                  [0,0,1,0,0,0,0,0],
                                  [0,0,0,1,0,0,0,0],
                                  [0,0,0,0,1,0,0,0],
                                  [0,0,0,0,0,1,0,0],
                                  [0,0,0,0,0,0,1,0],
                                  [0,0,0,0,0,0,0,1],
                                  [0,0,0,0,0,0,0,1]]), index=states,columns=states)
```

P_normal

```
##      0  10  20  30  40  50  60  70
## 0    0   1   0   0   0   0   0   0
## 10   0   0   1   0   0   0   0   0
## 20   0   0   0   1   0   0   0   0
## 30   0   0   0   0   1   0   0   0
## 40   0   0   0   0   0   1   0   0
## 50   0   0   0   0   0   0   1   0
## 60   0   0   0   0   0   0   0   1
## 70   0   0   0   0   0   0   0   1
```

```
P_speed=pd.DataFrame(np.matrix([[.1,0,.9,0,0,0,0,0],
                                  [.1,0,0,.9,0,0,0,0],
                                  [0,.1,0,0,.9,0,0,0],
                                  [0,0,.1,0,0,.9,0,0],
                                  [0,0,0,.1,0,0,.9,0],
                                  [0,0,0,0,.1,0,0,.9],
                                  [0,0,0,0,0,.1,0,.9],
                                  [0,0,0,0,0,0,0,1]]), index=states, columns=states)
```

P_speed

```
##      0  10  20  30  40  50  60  70
## 0    0.1 0.0 0.9 0.0 0.0 0.0 0.0 0.0
## 10   0.1 0.0 0.0 0.9 0.0 0.0 0.0 0.0
## 20   0.0 0.1 0.0 0.0 0.9 0.0 0.0 0.0
## 30   0.0 0.0 0.1 0.0 0.0 0.9 0.0 0.0
## 40   0.0 0.0 0.0 0.1 0.0 0.0 0.9 0.0
```

```
## 50  0.0  0.0  0.0  0.0  0.1  0.0  0.0  0.9
## 60  0.0  0.0  0.0  0.0  0.0  0.1  0.0  0.9
## 70  0.0  0.0  0.0  0.0  0.0  0.0  0.0  1.0
```

```
R_s_a=pd.DataFrame(np.c_[[-1,-1,-1,-1,0,-1,-1,0],[-1.5,-1.5,-1.5,-1.5,-0.5,-1.5,-1.5,0]],index=states,columns=
#np.c_[[-1,-1,-1,-1,0,-1,-1,0],[-1.5,-1.5,-1.5,-1.5,-0.5,-1.5,-1.5,0]],index=states,columns=['n','s'])
```

```
R_s_a.T
```

```
##      0   10   20   30   40   50   60   70
## n -1.0 -1.0 -1.0 -1.0  0.0 -1.0 -1.0  0.0
## s -1.5 -1.5 -1.5 -1.5 -0.5 -1.5 -1.5  0.0
```

```
pi_speed=pd.DataFrame(np.c_[np.repeat(0,len(states)),np.repeat(1,len(states))],index=states,columns=['n','s'])
```

```
pi_50=pd.DataFrame(np.c_[np.repeat(0.5,len(states)), np.repeat(0.5,len(states))],index=states,columns=['n','s'])
```

```
pi_speed.T
```

```
##      0   10   20   30   40   50   60   70
## n  0    0    0    0    0    0    0    0
## s  1    1    1    1    1    1    1    1
```

```
pi_50.T
```

```
##      0   10   20   30   40   50   60   70
## n  0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5
## s  0.5  0.5  0.5  0.5  0.5  0.5  0.5  0.5
```

Simulator - π^{speed}

```
pi=pi_speed
np.random.seed(1234)

history=list()
MC_N=10000

for MC_i in range(MC_N):
    s_now='0'
    history_i=list(s_now)

    while s_now != '70' :
        if np.random.uniform(0,1) < pi.loc[s_now]['n']:
            a_now='n'
            P=P_normal
        else:
            a_now='s'
            P=P_speed

        r_now=str(R_s_a.loc[s_now][a_now])
        s_next=states[np.argmin(P.loc[s_now].cumsum() < np.random.uniform(0,1))].item()
        history_i.extend([a_now,r_now,s_next])
        s_now=s_next

    history.append(history_i)

history_speed=history

pd.Series(map(lambda x: ','.join(x), history_speed[:20]))
```

```
## 0      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 1      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 2      0,s,-1.5,0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1...
## 3      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 4      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 5      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 6      0,s,-1.5,20,s,-1.5,10,s,-1.5,30,s,-1.5,50,s,-1...
## 7      0,s,-1.5,20,s,-1.5,40,s,-0.5,30,s,-1.5,50,s,-1...
## 8      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 9      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
```

```

## 10      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 11      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 12      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 13      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 14      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 15      0,s,-1.5,20,s,-1.5,10,s,-1.5,30,s,-1.5,50,s,-1...
## 16      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 17      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 18      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## 19      0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## dtype: object

```

Simulator - π^{50}

```
pi=pi_50
np.random.seed(1234)

history=list()
MC_N=10000

for MC_i in range(MC_N):
    s_now='0'
    history_i=list(s_now)

    while s_now != '70' :
        if np.random.uniform(0,1) < pi.loc[s_now]['n']:
            a_now='n'
            P=P_normal
        else:
            a_now='s'
            P=P_speed

        r_now=str(R_s_a.loc[s_now][a_now])
        s_next=states[np.argmin(P.loc[s_now].cumsum() < np.random.uniform(0,1))].item()
        history_i.extend([a_now,r_now,s_next])
        s_now=s_next

    history.append(history_i)

history_50=history

pd.Series(map(lambda x: ','.join(x), history_50[:20]))
```

```
## 0      0,n,-1.0,10,n,-1.0,20,s,-1.5,40,n,0.0,50,s,-1....
## 1      0,n,-1.0,10,s,-1.5,30,n,-1.0,40,s,-0.5,30,s,-1...
## 2      0,s,-1.5,20,n,-1.0,30,n,-1.0,40,s,-0.5,60,s,-1...
## 3      0,s,-1.5,20,n,-1.0,30,n,-1.0,40,n,0.0,50,n,-1....
## 4      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,s,-1.5,20,s,-1...
## 5      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,n,-1.0,40,n,0....
## 6      0,n,-1.0,10,n,-1.0,20,n,-1.0,30,n,-1.0,40,n,0....
## 7              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,n,-1.0,70
## 8      0,s,-1.5,20,n,-1.0,30,s,-1.5,50,n,-1.0,60,s,-1...
## 9      0,s,-1.5,20,s,-1.5,40,n,0.0,50,n,-1.0,60,n,-1....
```

```

## 10      0,n,-1.0,10,s,-1.5,30,n,-1.0,40,s,-0.5,60,s,-1...
## 11      0,s,-1.5,20,n,-1.0,30,n,-1.0,40,n,0.0,50,n,-1...
## 12      0,n,-1.0,10,s,-1.5,30,n,-1.0,40,n,0.0,50,s,-1...
## 13              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 14              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 15              0,n,-1.0,10,s,-1.5,30,s,-1.5,50,s,-1.5,70
## 16      0,s,-1.5,20,s,-1.5,40,n,0.0,50,n,-1.0,60,n,-1...
## 17              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,n,-1.0,70
## 18      0,n,-1.0,10,n,-1.0,20,s,-1.5,40,n,0.0,50,n,-1...
## 19              0,s,-1.5,20,s,-1.5,40,s,-0.5,60,s,-1.5,70
## dtype: object

```

Implementation 1 - π^{speed} (vectorized)

```
pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','sum'])
```

```
pol_eval.T
```

```
##           0   10   20   30   40   50   60   70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## sum    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(len(history_speed)):
    history_i=history_speed[MC_i]

    for j in range(0,len(history_i),3):
        pol_eval.loc[history_i[j]]['count']+=1

        if j < len(history_i) :
            pol_eval.loc[history_i[j]]['sum']+=np.sum(np.array(history_i[j+2:len(history_i)-1:3])).astype(float)
        else:
            pol_eval.loc[history_i[j]]['sum']+=0
```

```
pol_eval.T
```

```
##           0       10       20       30       40       50       60       70
## count 11225.0 1076.0 10291.0 1887.0  9485.0  2563.0  8563.0 10000.0
## sum   -65136.0 -5619.0 -42703.0 -6539.0 -22275.5 -4472.5 -14355.0      0.0
```

```
pol_eval['sum']/pol_eval['count']
```

```
## 0    -5.802762
## 10   -5.222119
## 20   -4.149548
## 30   -3.465289
## 40   -2.348498
## 50   -1.745025
## 60   -1.676398
## 70    0.000000
## dtype: float64
```


Implementation 2 - π^{speed} (running estimate)

```
pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','est'])
```

```
pol_eval.T
```

```
##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(len(history_speed)):
    history_i=history_speed[MC_i]

    for j in range(0,len(history_i),3):
        # update count
        pol_eval.loc[history_i[j]]['count']+=1
        current_cnt=pol_eval.loc[history_i[j]]['count']

        # return is the new info
        if j < len(history_i):
            new_info=np.sum(np.array(history_i[j+2:len(history_i)-1:3]).astype(float))

        else:
            new_info=0

        # update the last estimate with new info
        alpha=1/current_cnt
        pol_eval.loc[history_i[j]]['est']+=alpha*(new_info-pol_eval.loc[history_i[j]]['est'])

np.round(pol_eval.T,2)
```

```
##           0         10         20         30         40         50         60         70
## count 11225.0 1076.00 10291.00 1887.00 9485.00 2563.00 8563.00 10000.0
## est   -5.8   -5.22   -4.15   -3.47   -2.35   -1.75   -1.68    0.0
```

Implementation 3 - π^{50} (vectorized)

```
pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','sum'])
```

```
pol_eval.T
```

```
##           0   10   20   30   40   50   60   70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## sum    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(len(history_50)):
    history_i=history_50[MC_i]

    for j in range(0,len(history_i),3):
        pol_eval.loc[history_i[j]]['count']+=1

        if j < len(history_i) :
            pol_eval.loc[history_i[j]]['sum']+=np.sum(np.array(history_i[j+2:len(history_i)-1:3])).astype(float)

        else:
            pol_eval.loc[history_i[j]]['sum']+=0
```

```
pol_eval.T
```

```
##           0           10           20           30           40           50           60           70
## count 10863.0   5792.0   8140.0   7121.0   7549.0   7363.0   6991.0  10000.0
## sum  -64904.5 -29662.5 -33549.0 -24133.0 -15410.0 -14874.5 -9436.5      0.0
```

```
pol_eval['sum']/pol_eval['count']
```

```
## 0    -5.974823
## 10   -5.121288
## 20   -4.121499
## 30   -3.388990
## 40   -2.041330
## 50   -2.020168
## 60   -1.349807
## 70    0.000000
## dtype: float64
```

Implementation 4 - π^{50} (running estimate)

```
pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','est'])
```

```
pol_eval.T
```

```
##           0   10   20   30   40   50   60   70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for MC_i in range(len(history_50)):
    history_i=history_50[MC_i]

    for j in range(0,len(history_i),3):
        # increment count
        pol_eval.loc[history_i[j]]['count']+=1
        current_cnt=pol_eval.loc[history_i[j]]['count']

        # return is the new info
        if j < len(history_i):
            new_info=np.sum(np.array(history_i[j+2:len(history_i)-1:3]).astype(float))

        else:
            new_info=0

        # update the last estimate with new info
        alpha=1/current_cnt
        pol_eval.loc[history_i[j]]['est']+=alpha*(new_info-pol_eval.loc[history_i[j]]['est'])

np.round(pol_eval.T,2)
```

```
##           0       10       20       30       40       50       60       70
## count 10863.00  5792.00  8140.00  7121.00  7549.00  7363.00  6991.00 10000.0
## est   -5.97   -5.12   -4.12   -3.39   -2.04   -2.02   -1.35    0.0
```

TD - Implementation 5 - π^{speed}

```
pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','est'])
```

```
pol_eval.T
```

```
##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for episode_i in range(len(history_speed)):
    history_i=history_speed[episode_i]

    for j in range(0,len(history_i),3):

        # update count
        pol_eval.loc[history_i[j]]['count']+=1
        current_cnt=pol_eval.loc[history_i[j]]['count']

        # build TD target
        if j+3 < len(history_i):
            TD_tgt=float(history_i[j+2])+pol_eval.loc[history_i[j+3]]['est']
        else:
            TD_tgt=0

        # TD-updating

        alpha=1/current_cnt

        pol_eval.loc[history_i[j]]['est']+=alpha*(TD_tgt-pol_eval.loc[history_i[j]]['est'])

np.round(pol_eval.T,2)
```

```
##           0         10        20        30        40        50        60        70
## count 11225.00 1076.00 10291.00 1887.00 9485.00 2563.00 8563.00 10000.0
## est   -5.74   -5.19   -4.13   -3.48   -2.35   -1.75   -1.68    0.0
```

TD - Implementation 6 - π^{50}

```
pol_eval=pd.DataFrame(np.zeros((len(states),2)), index=states, columns=['count','est'])
```

```
pol_eval.T
```

```
##           0    10    20    30    40    50    60    70
## count  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
## est    0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0
```

```
for episode_i in range(len(history_50)):
    history_i=history_50[episode_i]

    for j in range(0,len(history_i),3):

        # update count
        pol_eval.loc[history_i[j]]['count']+=1
        current_cnt=pol_eval.loc[history_i[j]]['count']

        # build TD target
        if j+3 < len(history_i):
            TD_tgt=float(history_i[j+2])+pol_eval.loc[history_i[j+3]]['est']
        else:
            TD_tgt=0

        # TD-updating

        alpha=1/current_cnt

        pol_eval.loc[history_i[j]]['est']+=alpha*(TD_tgt-pol_eval.loc[history_i[j]]['est'])

np.round(pol_eval.T,2)
```

```
##           0          10          20          30          40          50          60          70
## count 10863.00  5792.00  8140.00  7121.00  7549.00  7363.00  6991.00 10000.0
## est   -5.84   -5.05   -4.08   -3.37   -2.04   -2.03   -1.35    0.0
```